

UNIVERSITY OF PENNSYLVANIA
ESE 650: LEARNING IN ROBOTICS
HOMEWORK 3

Changelog: This space will be used to note down updates/errata to the homework problems.

Read the following instructions carefully before beginning to work on the homework.

- You will submit neatly written solutions via Gradescope. You can use \LaTeX (encouraged) but you can also write by hand. You can use `hw_template.tex` on Canvas in the “hw” folder to do so. If your handwriting is *unambiguously legible*, you can submit PDF scans/tablet-created PDFs.
- Please start a new problem on a fresh page and mark all the pages corresponding to each problem. Failure to do so may result in your work not graded completely.
- Clearly indicate the name and Penn email ID on your submitted solutions.
- For each problem in the homework, you should mention the total amount of time you spent on it.
- You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel free to draw it on paper clearly, click a picture and include it in your solution. Do not spend undue time on typesetting solutions.
- You will see an entry of the form “HW 0 PDF” where you will upload the PDF of your solutions. You will also see entries like “HW 0 Problem 1 Code” where you will upload your solution for the respective problems. **For each programming problem, you should create a fresh Python file.** This file should contain all the code to reproduce the results of the problem and you will upload the `.py` file to Gradescope. If we have installed Autograder for a particular problem, you will use the Autograder. Name your file to be “`pennkey_hw0_problem1.py`”, e.g., I will name my code for Problem 1 as “`pratikac_hw0_problem1.py`”.
- **You should include all the relevant plots in the PDF, without doing so you will not get full credit.** You can, for instance, export your Jupyter

notebook as a PDF (you can also use text cells to write your solutions) and export the same notebook as a Python file to upload your code.

- **Your PDF solutions should be completely self-contained. If the question requires you to produce a plot, you must have it in the PDF to get credit. We will run the Python file to check if your solution reproduces the results in the PDF.**

Credit. The points for the problems add up to 115. You only need to solve for 100 points to get full credit, i.e., your final score will be $\min(\text{your total points}, 100)$.

1 **Problem 1 (Policy Iteration, 20 points (No Autograder)).** Consider the following
 2 Markov Decision Process. The state-space is a 10×10 grid, cells that are obstacles
 3 are marked in gray. The initial state of the robot is in blue and our desired terminal
 4 state is in green. The robot gets a *reward* of 10 if it reaches the desired terminal state
 5 with a discount factor of 0.9. At each non-obstacle cell, the robot can attempt to
 6 move to any of the immediate neighboring cells using one of the four controls (North,
 7 East, West and South). The robot cannot move diagonally. The move succeeds with
 8 probability 0.7 and with the remainder probability of 0.3, the robot can end up at
 9 some other cell. In short,

$$\begin{aligned} P(\text{moves north} \mid \text{control is north}) &= 0.7, \\ P(\text{moves west} \mid \text{control is north}) &= 0.1, \\ P(\text{moves east} \mid \text{control is north}) &= 0.1, \\ P(\text{does not move} \mid \text{control is north}) &= 0.1. \end{aligned}$$

10

11 Similarly, if the robot desired to go east, it may end up in the cells to its north, south,
 12 or stay put at the original cell with total probability 0.3 and actually move to the
 13 cell east with probability 0.7. The robot pays a cost of 1 (i.e., reward is -1) for each
 14 control input it takes, regardless of the outcome. If the robot ends up at a state
 15 marked as an obstacle (all grey cells are obstacles, i.e., cell marked (0,0), (0,1), (3,2)
 16 etc. are obstacles), it gets a reward of -10 for each time-step that it remains inside
 17 the obstacle cell. The robot is allowed to stay in the goal state indefinitely (i.e., take
 18 a special action to “not move”) and this action gets no reward/cost.

19 We would like to implement policy iteration to find the best trajectory for the
 20 robot to go from the blue cell to the green cell.

- 21 (a) **(0 points)** Carefully code up the above environment to run policy iteration.
 22 You will need to think about how to code up the probability transition matrix
 23 $\mathbb{R}^{100 \times 100} \ni T_{x,x'}(u) = P(x' \mid x, u)$, the run-time cost $q(x, u)$, and the
 24 terminal cost $q_f(x)$. Policy iteration is easy to implement if you represent
 25 all the above quantities as matrices and vectors. Plot the environment to
 26 check if it confirms to the above picture.
 27 (b) **(10 points)** Initialize policy iteration with a feedback control $u^{(0)}(x)$ where
 28 the robot always goes east, this results in a policy $\pi^{(0)} = (u^{(0)}(\cdot), u^{(0)}(\cdot), \dots)$.
 29 Write the code for policy evaluation to obtain the cost-to-go from every cell
 30 in the above picture for this initial policy. Plot the value function $J^{\pi^{(0)}}(x)$
 31 as a heatmap in the above picture.

32 (c) **(10 points)** Execute the policy iteration algorithm, you will iteratively
33 perform policy evaluation and policy improvement steps. For the first 4 iterations,
34 plot the feedback control $u^{(k)}(x)$ (using arrows as shown in the lecture
35 notes (https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.arrow.html,
36 you can also write the control input in the cell). You should color the cell
37 using the value function $J^{\pi^{(k)}}(x)$.

38 We have left the transition probabilities and the reward structure a bit vague to
39 force you to think carefully of the nuances of this problem. But some clarification
40 could be useful.

- 41 (1) You can code up what are called “sticky obstacles”, i.e., if the robot enters
42 an obstacle, then it stays there forever while incurring the obstacle cost at
43 each time instant.
- 44 (2) It is easiest to think of the runtime cost in this problem as a function of three
45 quantities $q(x, u, x')$ where x is the current state, u is the control and x' is
46 the next state. The Bellman equation then becomes

$$J^*(x) = \min_{u \in U} \mathbb{E}_{x'} [q(x, u, x') + \gamma J^*(x')] .$$

47 You will submit your own code for this problem, there is no Autograder.


48 **Problem 2 (Simultaneous Localization and Mapping (SLAM) with a particle**
49 **filter, 60 points (No Autograder).).**

50

51 **Download the data from the following link**

52 https://drive.google.com/file/d/1X92V8ISgV50KckK199v0BW_WGGRbQ1Ct/view?usp=sharing.

53 In this problem, we will implement mapping and localization in an outdoor
54 environment using information from a GPS/IMU system and a LiDAR sensor. We
55 have provided you with a subset of the KITTI visual odometry dataset collected
56 from a laserscanner mounted on a Volkswagen Passat B6. You can read more about
57 the hardware setup on the website (<https://www.cvlibs.net/datasets/kitti/setup.php>),
58 original paper (<https://www.cvlibs.net/publications/Geiger2012CVPR.pdf>), and the
59 video (https://www.youtube.com/watch?v=KXpZ6B1YB_k).




60

61 **Hardware setup of ‘Annieway’.** The autonomous driving platform (Annieway) has
62 a Velodyne HDL-64E laser scanner mounted on the roof, which is 1.73m above the
63 ground. The laser scanner has a vertical resolution of 64 and rotates at 10 frames
64 per second, capturing approximately 100,000 points per rotation. Additionally, the
65 autonomous driving platform is equipped with 2 color and grayscale cameras with
66 each color and grayscale camera mounted on the center and right side of the car.

67 The second type of observations we will utilize corresponds to the car’s location.
68 However, in contrast to the previous homework where we used the raw accelerometer
69 and gyroscope readings to get the orientation, we will use the (x, y, z) pose of

70 the center camera. These poses were created by the GPS/IMU system with RTK
 71 float/integer corrections (read up on wikipedia on what RTK is) enabled. We will
 72 transform the poses to coordinate system (x, z, θ) . Conceptually, these RTK poses
 73 are the same as Vicon in the previous homework, these are a precise estimate of the
 74 car's pose that we will use to check how well our SLAM system is working.

75 **Coordinate frames.** Poses are in coordinate system of the the left camera which
 76 has the Z axis pointing forwards, the X axis pointing right and the Y axis pointing
 77 downwards (see the picture below). However, the LiDAR has a coordinate frame
 78 with the X axis pointing forwards, Y axis pointing left and the Z axis pointing
 79 upwards. The camera is at a height of 1.65m and the height of the LiDAR is 1.73m
 80 above the ground. The world coordinate frame where we want to build the map has
 81 its origin on the X-Z ground plane.



83 Data and code.

85 (a) **(0 points)** We have provided you 4 datasets corresponding to 4 different
 86 trajectories of the car driving around the mid-size city of Karlsruhe, in rural areas
 87 and on highways. We have modified the original KITTI dataset by adjusting the
 88 LiDAR frequency from 10Hz to 2Hz (so that you do not have to process very
 89 large point clouds). The datasets are labeled as 00, 01, 02 and 03 in our modified
 90 version, they correspond to trajectories 02, 06, 07 and 08 from the original KITTI
 91 dataset, respectively. The provided dataset consists of three main directories: **calib**,
 92 **odometry** and **poses**.

93 The **calib** folder contains a .txt file that provides calibration data for the cameras.
 94 The P0/P1/P2/P3 are the 3×4 projection matrices after rectification, where P0/P2
 95 corresponds to the left cameras, and P1/P3 denotes the right cameras (refer to

96 figure above for details of the camera). The `Tr` transforms a point from Velodyne
97 coordinates into the left rectified camera coordinate system. Additionally, the
98 `times.txt` file stores timestamps for each image pair in seconds; however, this file will
99 not be used for this homework. We will not use any observations from the cameras
100 in this homework.

101 The **odometry** folder contains Velodyne point cloud data stored in a binary format
102 to save space. Each scan is represented as an $N \times 4$ floating-point matrix, where the
103 first three values correspond to the x , y , and z coordinates, and the fourth value
104 represents the reflectance. This data is organized in a row-aligned format, meaning
105 that the first four values in the file correspond to the first recorded measurement. The
106 function `load_kitti_lidar_data` in `load_data.py` is used to read the odometry
107 data, while the functions `show_kitti_lidar` and `show_kitti_lidar_sequence`
108 provide visualization of the KITTI dataset.

109 The **poses** folder contains trajectory poses for the entire sequence. Each `.txt` file
110 stores an $N \times 12$ table, where N represents the number of frames in the sequence.
111 Each i row corresponds to the i -th pose of the left camera coordinate system (with
112 the Z -axis pointing forward) represented by a 3×4 transformation matrix. The
113 transformation matrices are stored in a row-aligned format, meaning that the first
114 entries correspond to the first row of the matrix. These matrices transform a point
115 from the i -th coordinate system to the first (0th) coordinate system. Consequently,
116 the translational component (the 3×1 vector in the fourth column) represents the pose
117 of the left camera coordinate system in the i -th frame relative to the first (0th) frame.
118 The function `load_kitti_poses` in `load_data.py` is used to read the pose data,
119 while the function `trajectory2d` in `load_data.py` provides a visualization of
120 the trajectory.

121 You should read these functions carefully and check the values returned by them.

122 (b) (0 points) Next look at the `slam.py` file provided to you. Read the code for
123 the class `map_t` and `slam_t` and the comments provided in the code very carefully.
124 You are in charge of filling in the missing pieces marked as `TODO: XXXXXX`.
125 A suggested order for studying this code is as follows: `slam_t.read_data`,
126 `slam_t.init_particles`, `slam_t.lidar2world`, `map_t.__init__`,
127 `map_t.grid_cell_from_xz`. Next, the file `utils.py` contains a few standard
128 rigid-body transformations that you will need. You should pay attention to the
129 functions `smart_plus_2d` and `smart_minus_2d` that will be used to code up the
130 dynamics propagation step of the particle filter.

131 (c) (10 points, dynamics step) Next look at `main.py` which has two functions
132 `run_dynamics_step` and `run_observation_step` which act as test functions to

133 check if the particle filter and occupancy grid update has been updated correctly.
 134 The `run_dynamics` function plots the trajectory of the car (as given by its pose
 135 data). It also initializes 3 particles and plots all particles at different time-steps while
 136 performing the dynamics step with a very small dynamics noise; this is a very neat
 137 way of checking if dynamics propagation in the particle filter is working correctly.
 138 This function will create two plots, one for the odometry trajectory and one more for
 139 the particle trajectories, both these trajectories should match after you code up the
 140 dynamics function `slam_t.dynamics_step` correctly.

141 (d) **(20 points, observation step)** The function `run_observation_step` is used
 142 to perform the observation step of the particle filter to get an estimate of the location
 143 of the robot and updates to the occupancy grid using observations from the LiDAR.
 144 First read the comments for the function `slam_t.observation_step` carefully.
 145 We first discuss the particle filtering part.

146 We first discuss the particle filtering part.

- 147 (i) For each particle, assuming that the particle is indeed the true position of the
 148 car, transform the LiDAR scanned point clouds into the world coordinates
 149 using the `slam_t.lidar2world` function.
- 150 (ii) In order to compute the updated weights of the particle, we need to know
 151 the likelihood of LiDAR scans given the state (our current occupancy grid
 152 in the case of SLAM). We are going to use a simple model to do so

$$\log P(\text{LiDAR scan as if the car is at particle } p \mid m) = \sum_{ij \in O} m_{ij} \quad (1)$$

153 where O is the set of occupied cells as detected by the LiDAR scan assuming
 154 the robot is at particle p and m_{ij} is our current estimate of the binarized
 155 map (more on this below). In simple words, if the occupied cells as given
 156 by our LiDAR match the occupied cells in the binarized map created from
 157 the past observations, then we say the log-probability of particle p is large.

- 158 (iii) You will next implement the function `slam_t.update_weights` that takes
 159 the log-probability of each particle p , its previous weights, calculates the
 160 updated weights of the particles.
- 161 (iv) Typically, resampling step (`slam_t.stratified_resampling`) is per-
 162 formed only if the effective number of particles (as computed in
 163 `slam_t.resample_particles`) falls below a certain threshold (30% in
 164 the code). Implement resampling as we discussed in the lecture notes.

165 **Mapping.** We have a number of particles $p^i = (x^i, z^i, \theta^i)$ that together give an
 166 estimate of the distribution of the location of the car. For this homework, you will
 167 only use the particle with the largest weight to update the map although typically

168 we update the map using all particles. Our goal is simple: we want to increase the
169 `map_t.log_odds` array at cells that are recorded as obstacles by the LiDAR and
170 decrease the values in all other cells. You should add `slam_t.log_odds_occ` to all
171 occupied cells and add `slam_t.log_odds_free` from all cells in the map. It is also
172 a good idea to clip the `log_odds` to like between `[-slam_t.map.log_odds_max,`
173 `slam_t.map.log_odds_max]` to prevent increasingly large values in the `log_odds`
174 array. The array `slam_t.map.cells` is a binarized version of the map (which is
175 used above to calculate the observation likelihood).

176 Check the `run_observation_step` function after you have implemented the
177 observation step.

178 (e) Since the map is initialized to zero at the beginning of SLAM which results in
179 all observation log-likelihoods to be zero in [1], we need to do something special for
180 the first step. We will use the first entry in `slam_t.poses` to get an accurate pose
181 for the robot and use its corresponding LiDAR readings to initialize the occupancy
182 grid. You can do this easily by initializing the particle filter to have just one particle
183 and simply calling the `slam_t.observation_step` as shown in `main.py`.

184 (f) **(30 points)** You will now run the full SLAM algorithm that performs one
185 dynamics step and observation step at each iteration in the function `run_slam` in
186 `main.py`. Make sure to start SLAM only after the time when you have both LiDAR
187 scans and joint readings (the two arrays start at different times). For all 4 datasets,
188 you will plot the final binarized version of the map, (x, z) location of the particle
189 in the particle filter with the largest weight at each time-step and the odometry
190 trajectory (x, z) (in a different color); this counts for 10 points each.

191 **Some Notes.** This problem is much easier and shorter than it may seem. You should
192 go through these steps carefully and in the suggested order. You should make sure
193 that the results of the previous step are correct before proceeding. The two functions
194 in `main.py` to check the dynamics and observation step are very important to find
195 bugs. You do not need to implement more than 100 lines of code.

196 **Problem 3 (Building a NeRF, 40 points (No Autograder; You can implement**
197 **this on Google Colab if necessary)).** NeRF is a technique for mapping complex
198 scenes by optimizing an underlying continuous volumetric representation using
199 a sparse set of input views. NeRFs represent the scene using a fully connected
200 (non-convolutional) deep network. The input to the network is 5-dimensional
201 $x \in \text{SE}(3)$ (without the roll). This consists of the 3-dimensional location in
202 Euclidean space, and two viewing directions. Using this input, the neural network
203 inside the NeRF outputs volume density $\sigma(x)$ and a view-dependent color $c(x)$ at
204 that spatial location. In this problem, we will implement a simplified NeRF, which
205 only takes 3D Euclidean coordinates (as you can imagine, the pictures from such a
206 NeRF do not change depending upon the viewpoint and therefore they will not look
207 as natural). We will implement the simplest possible version of a NeRF without
208 a lot of bells and whistles that are used in actual implementations, on downsized
209 training images. This way, the model will be small enough to train locally on your
210 laptop, or on Google Colab.

211 **(a) Data Loading and COLMAP (5 points).** On Canvas, we provide a dataset
212 consisting of 100 LEGO images captured from various angles (you are also encour-
213 aged to capture your own dataset and show results on it). You will use COLMAP,
214 a Structure-from-Motion (SfM), and Multi-View Stereo (MVS) pipeline to obtain
215 camera extrinsic estimation. COLMAP is an open-source library that is compatible
216 with Mac (install using Homebrew), Linux, and Windows. Install COLMAP first
217 following instructions provided in the [COLMAP documentation](#) or the one that
218 [NeRF Studio](#) provides.

219 Assuming the images are taken by the same camera (images have the same
220 intrinsic parameters, i.e., the same camera calibration), you should use COLMAP to
221 reconstruct a sparse model. The package also comes with a GUI (you can call it
222 using “colmap gui”) that provides a great interface and visualization. After getting
223 the sparse model, you will have to understand the provided colmap2nerf script and
224 use it to transform the sparse model file into a JSON file which contains information
225 such as camera intrinsic and extrinsic corresponding to each image. We will need
226 this information to begin training the NeRF.

227 **To simplify this task, we have also provided the correct JSON file on Canvas.**
228 **You are welcome to use it if you do not want to run COLAMP to get the camera**
229 **intrinsics and extrinsics. You will not get these 5 points, but this way you will**
230 **be able to proceed with the rest of the problem.**

231 You will then implement the `load_colmap_data` function, which reads in the
232 generated JSON file as well as the raw images. We recommend you resize the raw
233 images to a lower resolution, for example, from 800×800 to 200×200 , so that it

234 is feasible to train everything on your laptop. After resizing the images, remember
235 to change the camera parameters (height, width, and focal length) accordingly. You
236 should report these parameters in the PDF and how you calculated them.

237 **(b) Implementation of the NeRF (20 points).** You will now implement four key
238 functions.

239 **The `get_rays` function.** Assuming a pinhole camera model, complete the `get_rays`
240 function, which takes camera intrinsic parameters (camera calibration matrix) and
241 extrinsic parameters (locations from where the images were collected) as input
242 and returns a set of rays in the world frame. Each ray starts from the camera origin
243 and passes through one of the pixels (see the figure in Section 4.5.1 in the lecture
244 notes). We will use the homogeneous coordinates. Given a point $x_c = (i, j, k, 1)$
245 in the camera frame, the point can be transformed from the camera frame to the
246 world frame with $x_w = T_w^c x_c$, where T_w^c denotes the 4×4 transformation matrix
247 obtained from the previous question.

248 It is useful to emphasize the coordinate convention. We will adhere to the standard
249 NeRF coordinate convention for camera coordinates: +X is right, +Y is up, and
250 +Z points back and away from the camera, i.e., the -Z direction corresponds to the
251 direction at which the looking at. It is important to note that other code-bases on the
252 Internet may adopt the COLMAP/OpenCV convention, where the Y and Z axes are
253 flipped compared to ours, but the +X axis remains the same. The world coordinate
254 system is oriented such that the up vector is +Z. The XY plane is parallel to the
255 ground plane.

256 **The `sample_points_from_rays` function.** Given a set of rays emanating from the
257 camera center, we will discretize each ray into segments to approximate the integrals
258 during volume rendering. Implement the `sample_points_from_rays` function, which
259 returns an array of N_{sample} points along each ray in world coordinates.

- 260 (a) With a rough estimate of the distance from the object to the camera, we can
261 determine the clipping thresholds s_{near} (the distance of the nearest point
262 of interest) and s_{far} (the distance of the farthest point of interest). Each
263 ray will only be evaluated within the range of s_{near} and s_{far} , which defines
264 the volume of interest. Given a fixed number of points N_{sample} , a small
265 $s_{\text{far}} - s_{\text{near}}$ means that sampled points along the ray are closer to each other;
266 this leads to a better estimation for the integral.
- 267 (b) You can sample uniformly along the ray. For enhanced performance,
268 consider incorporating some randomness into the sampling process while
269 ensuring that there is at least one point every $(s_{\text{far}} - s_{\text{near}})/N_{\text{sample}}$.

270 **The position_encoding function.** Like we discussed in the lecture, an MLP with a
 271 finite width and a certain number of layers may not be able to represent functions
 272 of arbitrarily high bandwidths which are necessary to get high-frequency textures.
 273 This leads to blurry images from the NeRF. A neat solution to this issue is to use a
 274 different representation for the inputs $x \in \mathbb{R}^3$. Instead of using x we use

$$\varphi(x) = (\sin(2^k x_1), \sin(2^k x_2), \sin(2^k x_3), \dots)_{k=0, \dots, 10}$$

275 where k is the frequency and we choose, say 10 different frequencies. The input
 276 layer of the MLP would therefore be 30- instead of 3-dimensional.

277 **The volume_rendering function.** Here you will implement the volume rendering
 278 function in Equation 4.31 in the lecture notes. In summary, given a ray with points
 279 at distances

$$s_i = s_{\text{near}} + \frac{i}{N_{\text{sample}}}(s_{\text{far}} - s_{\text{near}})$$

280 we will calculate

$$\text{opacity: } \alpha_j = 1 - e^{-\sigma(s_j)(s_{j+1} - s_j)}$$

$$\text{transmittance: } p(s_i) = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

$$\text{color: } c = \sum_{i=1}^{N_{\text{sample}}} c(s_i) \alpha_i \prod_{j=1}^{i-1} (1 - \alpha_j).$$

281 for each ray corresponding to each pixel. Implement the volume_rendering function,
 282 which renders an RGB image using the predicted radiance field.

283 **(d) Network Training (10 points).** We provide the neural network architecture
 284 and a simple training loop for you to start. Fill in the **nerf_step_forward** function
 285 with your implementations of the functions above. Your report should mention
 286 the parameters for the **train** function, including s_{near} , s_{far} , and N_{sample} . Start with
 287 N_{sample} of 32 and the hidden dimension of the MLP h_{dim} of 32. With this setting,
 288 you should be able to train the network on a laptop CPU in about 20 minutes.

289 (i) We highly recommend rendering and visualizing the network prediction
 290 every few iterations (doing so is similar to calculating the validation loss
 291 after few epochs while training a standard neural network-based classifier).
 292 This is an easy way to assess the network's performance. You can select one
 293 of the poses from the training set or randomly select your own pose as the test
 294 pose. Then, render an RGB image at the test pose using **nerf_step_forward**
 295 and check if the rendered image makes sense.

296 (ii) If you have access to a GPU, or decide to use Colab, consider increasing
297 N_{sample} and h_{dim} . Doing so should lead to improved results.

298 You should report a plot of the training loss as a function of the number of weight
299 updates. You should report the final training loss, and for about 5-6 randomly
300 sampled images from the training dataset, you should show the original image
301 and the one rendered from the NeRF from the same viewpoint (this is reporting
302 predictions of the network on the training samples).

303 **(e) Inference (5 points).** Take the trained network, randomly pick 5 viewpoints
304 from as different poses as you can and report the rendered RGB images from these
305 viewpoints. Try to find viewpoints where the NeRF is working well as well as ones
306 where it is not.